EVALUATE THE DASK DISTRIBUTED COMPUTING FRAMEWORK IN RESPECT TO VARIOUS SCIENTIFIC COMPUTING TASKS

EERO VAINIKKO

Course Coordinator

ARTJOM LIND

Topic Supervisor

JEYHUN ABBASOV

Student

The Research Questions

How DASK's parallel and out-of-core computation extends the effective scale of modern hardware to larger datasets?

How these ideas can be more broadly applied to other parallel collections?

How the DASK performs on various scientific computing tasks?

- such as Average Global Ocean Temperature 36 year's worth

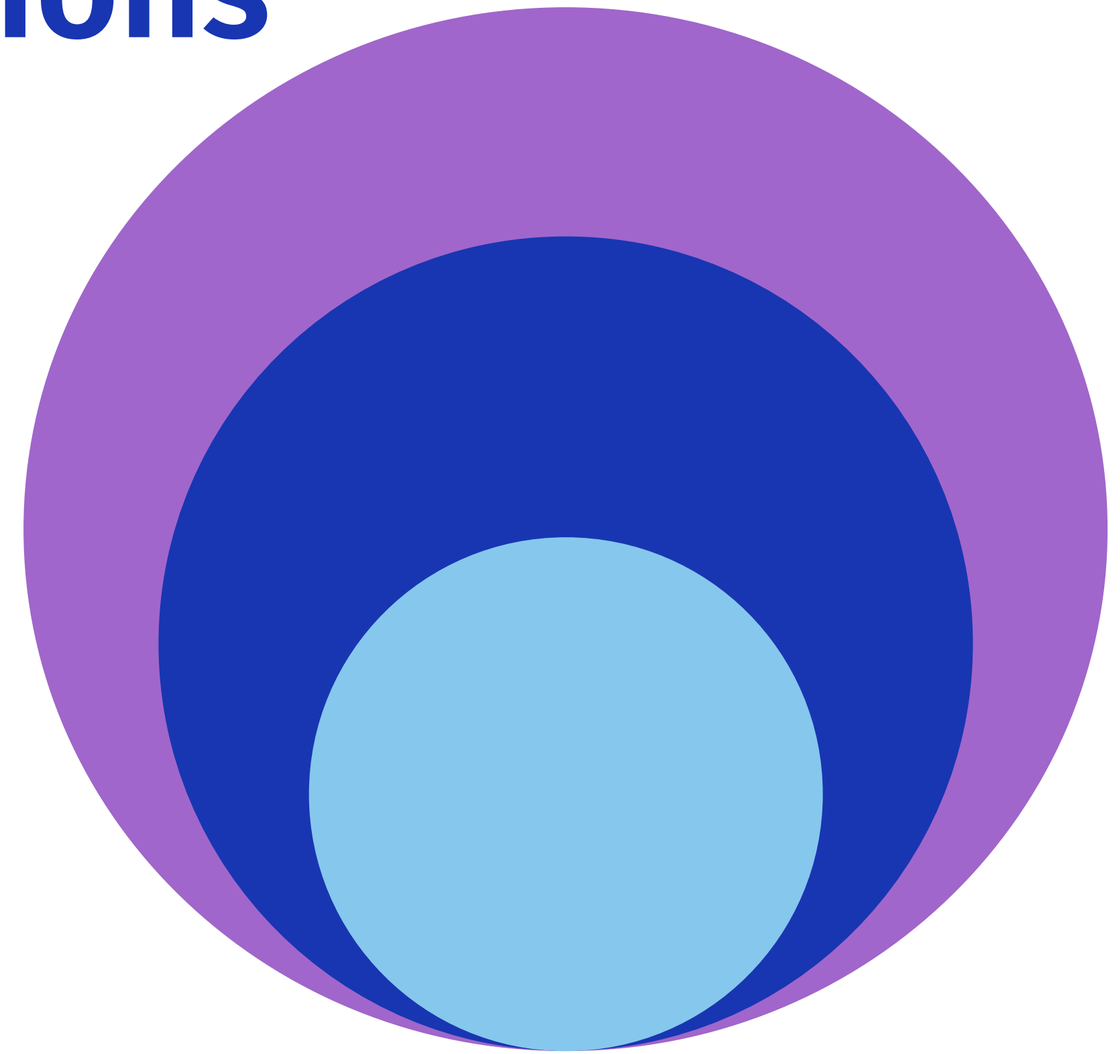

About Dask

A flexible library for parallel computing in Python.

Is composed two parts:

1

Dynamic task scheduling

- optimized for computation
- ***similar*** to Apache Airflow, Luigi Workflow, ***but*** optimized for interactive computational workloads

2

“Big Data” collections

- contains parallel arrays, dataframes, and lists that extend common interfaces like NumPy, Pandas, or Python iterators to larger-than-memory or distributed environments

Dask Architecture

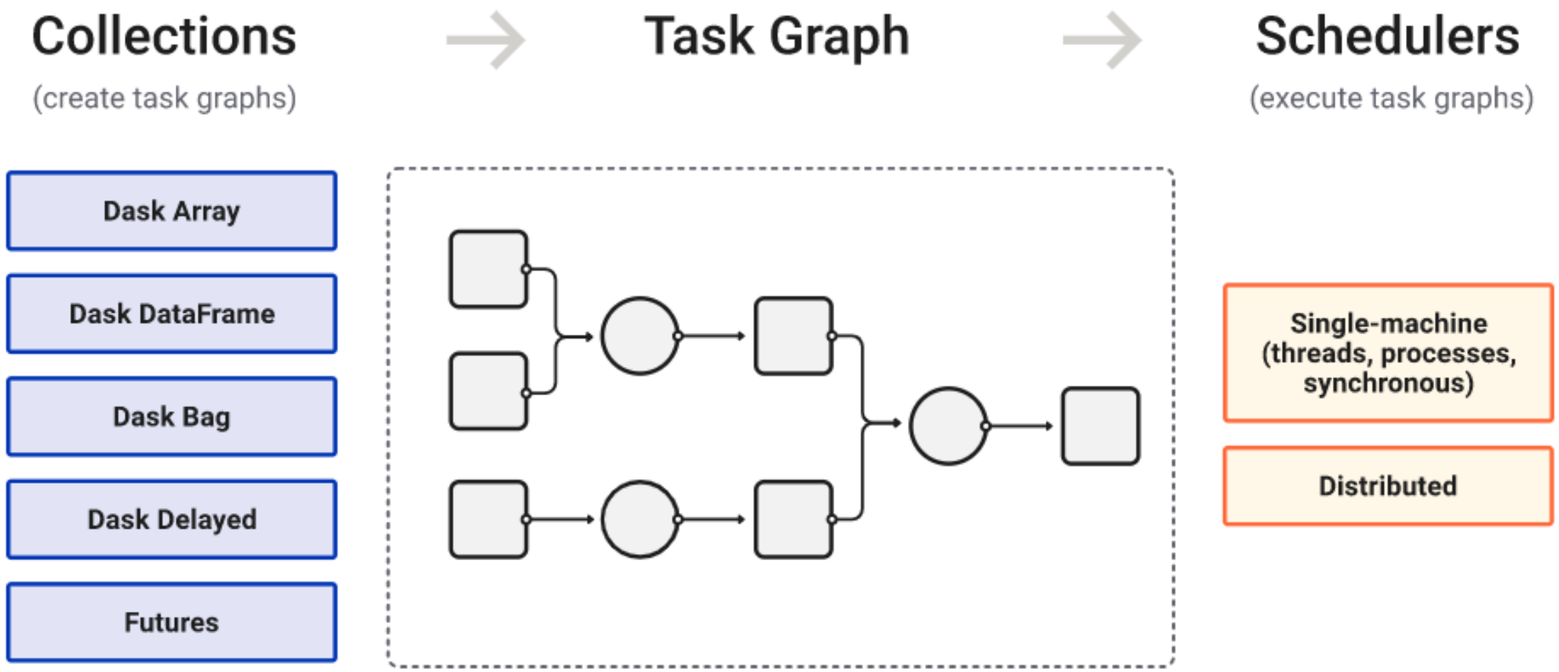

High level collections are used to generate task graphs which can be executed by schedulers on a single machine or a cluster.

Source: https://docs.dask.org/en/stable/

References

[1] Dask: Parallel Computation with Blocked algorithms and Task Scheduling

https://www.researchgate.net/publication/328778461_Dask_Parallel_Computation_with_Blocked_algorithms_and_Task_Scheduling

[2] Efficient MPI-based Communication for GPU-Accelerated Dask Applications

https://ieeexplore.ieee.org/abstract/document/9499534

[3] Performance Evaluation of Python Based Data Analytics Frameworks in Summit: Early Experiences

https://link.springer.com/chapter/10.1007/978-3-030-63393-6_24

[4] A Performance Comparison of Dask and Apache Spark for Data-Intensive Neuroimaging Pipelines

https://ieeexplore.ieee.org/abstract/document/8943502